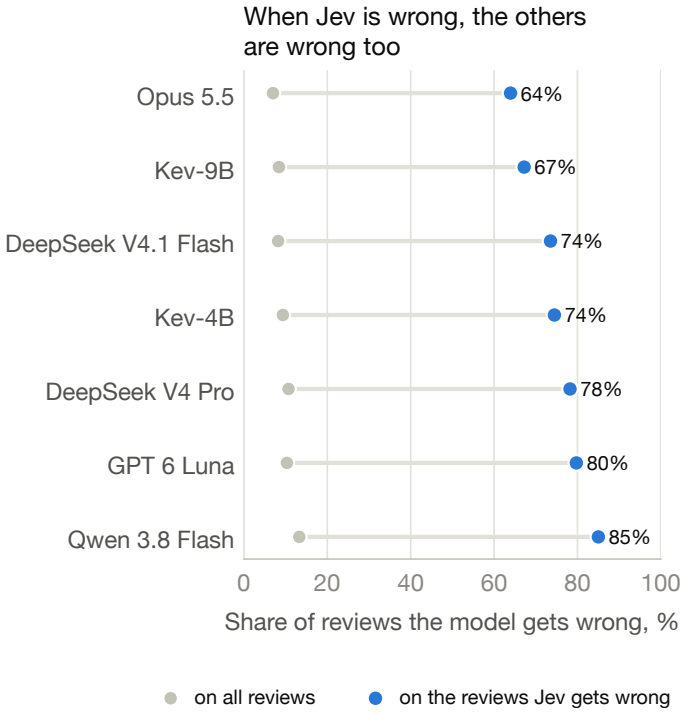
When Jev is wrong, the others are wrong too
Opus 5.5
Kev-9B
DeepSeek V4.1 Flash
Kev-4B
DeepSeek V4 Pro
GPT 6 Luna
Qwen 3.8 Flash
64%
67%
74%
74%
78%
80%
85%
0
20
40
60
80
100
Share of reviews the model gets wrong, %
on all reviews
on the reviews Jev gets wrong

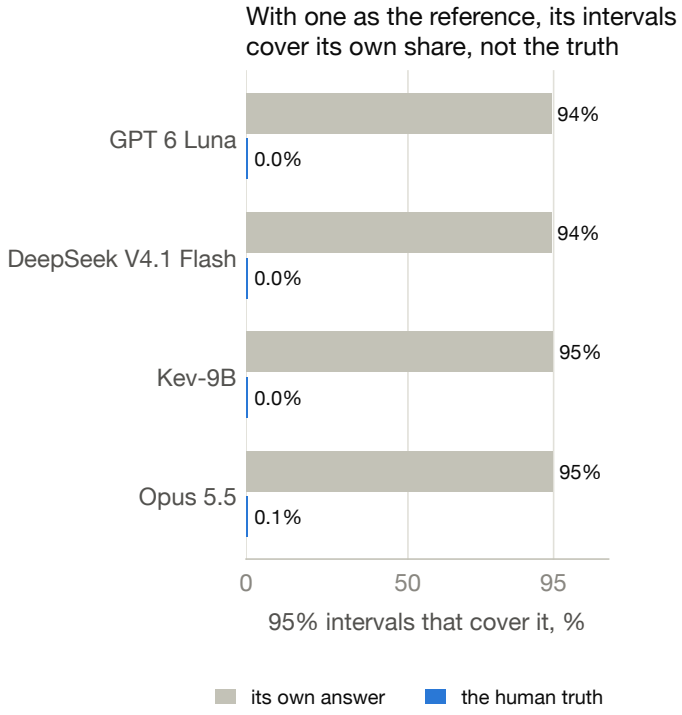
With one as the reference, its intervals cover its own share, not the truth
GPT 6 Luna
DeepSeek V4.1 Flash
Kev-9B
Opus 5.5
94%
0.0%
94%
0.0%
95%
0.0%
95%
0.1%
0
50
95
95% intervals that cover it, %
its own answer
the human truth